

Recent research on adaptive mechanisms in RNNs

"Identify recent peer-reviewed research on adaptive temporal scaling, adaptive update rates, time-scale mechanisms, and context-dependent recurrent state updates in RNNs for time-series forecasting. Compare these approaches with FAC-RNN and identify what remains missing." *Generated by LeapSpace*

Adaptive Recurrent Forecasting at a Glance

Key Findings Table

Topic	What recent peer-reviewed work shows	Populations / data settings	Implication for FAC-RNN comparison	Citations
Adaptive temporal scaling	Recent forecasting research increasingly replaces single fixed recurrence rates with multi-scale routing, adaptive scale weighting, and cross-scale interaction mechanisms.	Multiscale, long-horizon, multivariate forecasting settings	Suggests an important capability class beyond plain attention/context fusion: explicit temporal-resolution control.	[1] [2] [3] [4] [5]
Context-aware scale choice	Some models choose scales from local content rather than fixed hyperparameters, using routing, attention, or learnable gating over scales.	Context-varying sequences; multiresolution temporal patterns	Points to a likely gap for FAC-RNN unless it explicitly modulates timescale per step or per regime.	[6] [7] [8] [9] [10] [11]
Adaptive update / refresh control	Selective gates, memory-control mechanisms, and segment-wise recurrence can improve information flow, reduce parameters, or reduce recurrent cost.	Long-horizon and efficiency-sensitive forecasting	Adds a second capability axis beyond feature attention: deciding when and how much to update recurrent state.	[12] [13] [14] [15] [16] [17]
Long-horizon weakness remains	Even gated or adaptive recurrent models still show memory decay, compounding error, and instability under long horizons, anomalies, or chaotic dynamics.	Long-sequence, seasonal, nonstationary, chaotic forecasting	Adaptive updates help, but do not solve robustness and stability.	[18] [19] [20] [21] [22] [23]
Irregular-time forecasting	For sparse, missing, or irregularly sampled series, continuous-time and gap-aware recurrent formulations are a major direction.	ICU, sensor, satellite, crop, multisensor data	This is a likely missing or at least unverified FAC-RNN capability.	[24] [25] [26] [27] [28] [29] [30]
FAC-RNN evidence status	No retrieved paper explicitly documents FAC-RNN architecture, training setup, or benchmark results.	N/A	Comparison can only be indirect via related attention/context-fusion RNNs.	[31] [32] [33]

Topic	What recent peer-reviewed work shows	Populations / data settings	Implication for FAC-RNN comparison	Citations
What remains missing	The literature still lacks strong evidence for FAC-RNN on adaptive per-step update rates, irregular timing, calibration, exogenous-variable robustness under shift, and same-dataset comparisons against stronger alternatives.	Real-world deployment conditions	These are the clearest unmet research needs.	[34] [35] [36] [37] [38]

Direct Answer

Recent peer-reviewed research on RNN forecasting clusters into three relevant lines. First, models increasingly adapt **temporal scale** using multi-scale routing, dynamic scale attention, or scale-conditioned aggregation instead of a single fixed recurrent timescale [1] [2] [6] [7] [8]. Second, a separate line adapts **state refresh and forgetting** through selective gating, memory-control mechanisms, and segment-wise recurrence, aiming to regulate how often hidden states are updated and how much prior memory is retained [12] [13] [14] [15] [17]. Third, irregular-time forecasting increasingly uses **gap-aware or continuous-time recurrence**, where elapsed time between observations directly shapes state evolution [25] [26] [27] [28] [29].

Compared with FAC-RNN, the supplied evidence supports only an **indirect** comparison, because no retrieved peer-reviewed paper explicitly identifies FAC-RNN’s architecture, training procedure, or dataset results [31] [32] [33]. Based on related attention/context-fusion recurrent models, FAC-RNN-like designs likely emphasize feature selection and context aggregation, but the adaptive-update literature adds capabilities that are either missing or unverified for FAC-RNN: explicit per-step timescale control, learned update-rate modulation, segment-wise memory refresh, and irregular-sampling handling [6] [12] [17] [25].

What remains missing most clearly is evidence that FAC-RNN supports or has been tested on: (1) context-dependent recurrent state refresh inside the cell, (2) long-horizon stability under nonstationarity and regime shifts, (3) sparse or irregularly sampled forecasting, (4) calibrated probabilistic forecasting, and (5) systematic exogenous-covariate evaluation under distribution shift [18] [21] [34] [36] [37]. Taken together, these findings suggest the strongest next step is a hybrid design that combines FAC-style attention/context fusion with explicit adaptive temporal control and continuous-time or gap-aware state evolution [8] [12] [25].

Confidence level: High - The broad directional answer is directly supported by multiple convergent sources, but the FAC-RNN comparison itself is necessarily indirect because primary FAC-RNN evidence is absent [1] [12] [25] [31].

Study Scope

- **Time period emphasized:** mostly 2019-2025, with some adjacent foundational recurrent forecasting evidence where needed [14] [16] [24].

- **Disciplines/settings covered:** multivariate forecasting, long-horizon forecasting, clinical irregular series, satellite/sensor forecasting, chaotic systems, and related attention-augmented recurrent prediction tasks [25] [27] [39] [40].
- **Method families covered:** RNN, LSTM, GRU, modified recurrent cells, segment-wise recurrent models, attention-augmented recurrent models, and continuous-time recurrent variants [16] [17] [29] [41].

Assumptions & Limitations

- The main limitation is evidential: no retrieved paper explicitly documents FAC-RNN, so all FAC-RNN comparisons are inferential and based on related feature-attention/context-fusion recurrent models rather than direct same-model evidence [31] [32] [33].
- Benchmark evidence for explicit cell-level adaptive temporal scaling on standard public forecasting sets such as Electricity, Traffic, Exchange Rate, M4, or M5 is sparse in the supplied material [16] [42] [43].
- Across the literature, robustness, calibration, and regime-shift evaluation are less consistently reported than point-forecast error metrics [34] [36] [37] [44].

Suggested Further Research

- Build a FAC-style recurrent model with an explicit learned update-rate controller or selective refresh gate [12] [15].
- Test FAC-style context fusion on irregularly sampled data with time-gap-aware or continuous-time recurrence [25] [26] [27].
- Evaluate calibration and robustness under distribution shift, not only RMSE/MAE [34] [35] [36].
- Run same-dataset comparisons against Transformers, reservoir computing, and hybrid baselines on long-horizon and chaotic settings [38] [40] [45].
- Standardize reporting of efficiency, memory, and latency alongside accuracy [17] [46] [47].

Adaptive Recurrent Forecasting: Full Report

Foundational Context

Three mechanisms matter when discussing “adaptive” recurrent forecasting.

First, **temporal scaling** asks at what resolution the model should summarize the past. A fixed-timescale RNN updates once per step with a single hidden-state dynamic, but many real series contain daily, weekly, seasonal, bursty, and event-driven structure simultaneously, making a single rate restrictive [3] [4] [48]. Recent work therefore adds multiple scales, hierarchical structure, or dynamic scale weighting to better match multiscale dynamics [1] [2] [5].

Second, **adaptive update rate** concerns how often and how strongly the hidden state should refresh. In standard GRU/LSTM-style recurrence, update and forget behavior is gated, but several recent papers treat this as an explicit design target, modifying cell structure, adding selective gates, or refreshing state segment-wise to improve long-horizon learning and efficiency [12] [13] [14] [17].

Third, **irregular-time recurrence** extends adaptivity from “which scale?” or “how much update?” to “what does elapsed time mean?” For missing or asynchronous observations, continuous-time and time-gap-aware models evolve latent state between observations rather than assuming equally spaced updates [24] [25] [29].

These distinctions are important for comparing the broader literature with FAC-RNN-like models, because attention/context fusion mainly addresses **what information is salient**, whereas adaptive recurrent mechanisms additionally address **when, at what temporal granularity, and under what elapsed-time dynamics** state updates should occur [8] [25] [41].

Confidence level: High - This framing is a synthesis directly grounded in the recurring mechanisms identified across the supplied evidence [1] [12] [24].

1. Adaptive Temporal Scaling in Recurrent Forecasting

1.1 Multi-timescale and scale-adaptive recurrent designs

A consistent pattern in the recent literature is that RNN-family forecasting improves when models explicitly represent more than one temporal scale and allow interaction across scales [3] [4] [5]. Papers combining decomposition, attention, and message passing report better capture of both short- and long-range dependencies than plain recurrent baselines [3] [4] [49]. More broadly, conventional RNNs are described as weak on multiscale, intermittent, and rapidly changing series, motivating hierarchical or scale-separated designs [48] [50] [51].

The literature also emphasizes that temporal adaptation is not just architectural ornamentation. Adaptive scale mechanisms are motivated by the fact that the best temporal granularity can vary by regime, variable, and forecast horizon [1] [2] [13]. This is a substantive departure from fixed hyperparameter choices such as a single window or uniform update granularity.

1.2 Context-aware scale selection versus fixed hyperparameters

The strongest examples of content-conditioned scale choice in the supplied evidence are not always classic RNN cells, but they are directly relevant to recurrent forecasting design. MFformer uses meta-learning-based routing and adaptive multi-scale modeling to select and aggregate only a subset of temporal scales rather than committing to one scale globally [6]. CGFB-GNN’s AMSTCM dynamically selects relevant scales via attention instead of relying on predefined scales, with effective scaling varying with the data state [7]. TAPformer likewise contrasts with static patch sizes by using a dynamic dual-scale attention mechanism whose local/global blending is learnable and content-aware [8].

Taken together, these findings suggest that the field is moving from **architectures with multiple available scales** to **architectures that learn which scale is active in context** [6] [7] [8]. That distinction matters for FAC-RNN comparison: feature/context fusion alone does not establish that temporal scale itself is being selected adaptively.

1.3 Benchmark evidence and limitations

The evidence base is notably thinner when the standard is narrowed to peer-reviewed 2019-2025 RNN papers with **explicit cell-level adaptive temporal scaling** and results on canonical benchmarks such as Electricity, Traffic, Exchange Rate, M4, or M5. The supplied material does not include a study clearly satisfying all of those conditions [16] [42] [43]. The closest example, NOA-LSTM, modifies LSTM components and reports improved error and training time on Jena weather and Appliance Energy Usage, but not on the requested standard multivariate benchmarks [16].

This means the trend toward adaptive temporal handling is real, but benchmark proof remains fragmented [3] [4] [16]. That fragmentation limits strong claims about broad superiority over fixed-timescale recurrent baselines.

Confidence level: Medium - The directional trend is well supported, but benchmark coverage for explicit cell-level adaptive scaling is sparse in the supplied evidence [1] [6] [16].

2. Context-Dependent Update Rates and Hidden-State Refresh

2.1 Adaptive gating and selective update control

Several papers directly target hidden-state update control rather than only adding external attention. The general claim across this group is that adaptive gating improves long-term prediction by regulating forgetting and information flow more precisely [12] [13] [14]. KOSLM is especially relevant: it embeds a Kalman-optimal selective mechanism within a state-space reformulation of LSTM and reports lower MSE than state-of-the-art methods on long-horizon datasets, framing selective propagation and forgetting as central to recurrent forecasting quality [12]. Selective RNNs with Random Connectivity Gated Units dynamically adjust model importance across LSTM-, GRU-, and MGU-type variants while reducing parameter count by about two-thirds and improving benchmark performance over training the cells separately [15].

NOA-LSTM provides a complementary message: altering internal LSTM components can improve both error and training time relative to a standard LSTM-style format [16]. A theory-oriented 2021 paper further argues that forget-gate design determines retention timescale and proposes recurrent structures representing longer timescales than conventional models [14].

These studies collectively frame recurrent update rate not as a fixed property of the cell, but as a controllable design dimension [12] [14] [15].

2.2 Segment-wise reset and recurrent refreshing

Another route to adaptive refresh is not per-step gate redesign, but **segment-wise recurrence**. SegRNN divides long sequences into continuous segments and models them at the segment level, reducing recurrent steps while preserving temporal continuity [17]. This design is motivated both by efficiency and by the need to control error accumulation over long histories [17] [46] [52]. Standard GRU reset gates already implement a basic memory-reset policy, but segment-wise designs make refresh timing a higher-level architectural decision [53] [54].

Residual forecasting strategies also address accumulation by predicting changes rather than absolute values, which can improve robustness under nonstationarity [55]. Benchmark snippets suggest compact recurrent models can remain competitive or superior in both efficiency and accuracy on some tasks [56] [57] [58].

2.3 Failure modes under long horizons and nonstationarity

Despite these mechanisms, the evidence repeatedly notes unresolved long-horizon weaknesses. Supplied studies report memory decay, vanishing or exploding gradients, and growing forecast error as horizon length increases, even for LSTM-family models [18] [19] [20]. Seasonal SST forecasting work reports better seasonal/interannual modeling than traditional methods, yet still identifies low long-term stability, weak multiscale representation, and poorer generalization under anomalies or long horizons [21]. More broadly, decomposition-based or multi-scale designs can still fail when trend/seasonality assumptions break down or data are weakly decomposable and nonstationary [59] [60].

So adaptive update control helps, but it does not fully solve stability under shift-heavy or difficult long-horizon forecasting.

Confidence level: High - Multiple sources converge on both the benefits of adaptive update control and the persistence of long-horizon failure modes [12] [17] [20] [21].

3. Irregularly Sampled, Sparse, and Event-Driven Forecasting

For irregular or missing-observation forecasting, the literature broadens the notion of adaptive recurrence. Reviews identify two main strategies: impute first, or modify gated RNNs to directly incorporate missingness, time intervals, and irregularity [24]. The second strategy is especially relevant to the user's query because it makes recurrence explicitly conditional on elapsed time.

Time-continuous recurrent models evolve hidden state between observations and have been used in ICU blood-glucose forecasting on electronic medical record and simulated data [25]. RGRU-D uses masking, time interval, bursty, and cumulative missing-rate variables on clinical ICU data, showing how a GRU-family forecaster can ingest missingness structure directly [26]. In Sentinel-1 displacement forecasting, adding time intervals as a feature helped TG-LSTM on one-step prediction, though TFT performed best for multi-step forecasting, showing that recurrent gap-awareness is useful but not always dominant [27]. ODE-RNNs on irregularly spaced crop-classification image sequences with cloud-driven missing observations outperformed LSTM, GRU, TCN, and Transformer baselines [28].

A broader methodological theme emerges: irregular forecasting increasingly relies on continuous-time latent dynamics, neural ODEs, neural flows, or controlled differential equation approaches rather than only discrete adaptive gates [29] [30]. This is an important contrast with FAC-RNN-style attention/context fusion, because the latter does not by itself address elapsed-time semantics.

The literature also clarifies trade-offs. On chaotic spatiotemporal benchmarks, reservoir computing can outperform BPTT-trained gated RNNs in predictive performance, long-term statistics, and training time when full-state data are available [40]. But with reduced-order data, RC can become unstable and diverge more easily than trained gated RNNs [40]. Thus recurrent approaches remain competitive in some sparse or partial-observation settings even when alternative continuous-time methods are attractive.

Confidence level: High - The shift toward gap-aware and continuous-time recurrent modeling is directly and repeatedly supported in the supplied evidence [25] [26] [28] [29].

4. FAC-RNN and Related Attention-Context Fusion Forecasters

4.1 What can be said directly

The critical evidential fact is simple: no retrieved paper explicitly describes FAC-RNN's exact architecture, training setup, or benchmark wins [31] [32] [33]. Therefore, the supplied evidence does **not** support direct claims about FAC-RNN's cell design, update rule, or empirical standing on standard forecasting datasets.

4.2 Closest related evidence

What the evidence does include are related attention/fusion recurrent models. DA-RNN combines recurrent encoder-decoder updates with input attention and temporal attention for multivariate forecasting [39] [61]. Other attention-recurrent hybrids are reported to improve forecasting accuracy or trend consistency in particular applications [62]. Ablation evidence is especially informative: attention substantially improves a vanilla RNN, reducing MAE and RMSE markedly, but adds less when the backbone is already an LSTM, implying that recurrent gating and external attention partly substitute for one another [41]. TCN-QV ablations similarly show that attention improves accuracy and stability when layered on top of temporal context modeling [63].

This related evidence implies that FAC-RNN-like models likely derive their gains primarily from **feature selection and context aggregation**, not necessarily from adaptive state-refresh logic [41] [63]. That is plausible, but it remains an inference rather than a direct FAC-RNN result.

4.3 Trade-offs versus lighter or alternative recurrent designs

The broader benchmark literature also shows that attention-augmented recurrent models are not uniformly dominant. Some lighter or simpler recurrent designs achieve stable gains with fewer parameters or lower computational cost, such as SG-RNN or MRN [47] [64]. Other studies report dataset-dependent gains or no statistically significant differences among several recurrent hybrids [65]. Meanwhile, Transformers are often noted as computationally expensive, while RNNs retain efficiency appeal in some settings [66] [67] [68].

Confidence level: Medium - The indirect positioning of FAC-RNN relative to related attention-context models is supported, but only indirectly because FAC-RNN itself is absent from the retrieved primary evidence [31] [41] [47].

5. What Remains Missing Relative to FAC-RNN

5.1 Missing verified capabilities

The clearest gap is the absence of direct evidence that FAC-RNN supports **adaptive per-step temporal control**. No supplied paper shows FAC-RNN selecting scales dynamically, modulating update rates, or refreshing hidden state segment-wise [31] [32] [33]. In contrast, the broader literature has multiple examples of these capabilities in adjacent recurrent designs [6] [12] [17].

A second likely missing capability is **irregular-time handling**. Continuous-time and gap-aware recurrent models explicitly encode elapsed time, missingness, or asynchronous observation patterns [25] [26] [27] [28]. No retrieved FAC-RNN evidence demonstrates anything comparable [31].

5.2 Evaluation gaps: calibration, exogenous variables, distribution shift

The supplied material indicates that recurrent architectures are convenient for exogenous-variable inclusion in principle, but actual benchmark evidence on exogenous-covariate robustness is limited in the excerpts [69]. Probabilistic-forecasting papers in the retrieved set do not clearly report calibration metrics or distribution-shift evaluation for RNN baselines [34] [35] [36]. No retrieved study explicitly states whether FAC-RNN-style models are weaker on these dimensions [34] [36] [69].

Accordingly, what is missing is not proof that FAC-RNN fails on these dimensions, but proof that it has been **tested** on them.

5.3 Comparative gaps against stronger alternatives

The final gap is comparative. The evidence notes that statistical models may remain stronger on robustness, efficiency, and automation for non-expert users, and that hybrid or reservoir-computing alternatives can outperform fully trained RNNs in some regimes [37] [38] [40] [45]. Yet no provided paper tests FAC-RNN against these alternatives on the same datasets with calibration, regime-shift robustness, and efficiency all reported together [37] [38].

Taken together, these findings suggest that the central missing contribution is not another attention variant, but an integrated recurrent forecasting study combining:

1. FAC-style feature/context fusion,
2. adaptive timescale or update-rate control,
3. gap-aware or continuous-time handling,
4. robustness and calibration evaluation,
5. strong same-dataset baselines including Transformers, RC, and hybrids [8] [12] [25] [37] [38].

Confidence level: High - These missing elements follow directly from repeated evidence gaps and the absence of explicit FAC-RNN documentation [31] [36] [37].

Conclusion

The recent peer-reviewed literature shows a real shift from fixed-rate recurrent forecasting toward **context-adaptive temporal control**, expressed through multi-scale selection, selective forgetting, segment-wise state refresh, and continuous-time hidden-state evolution [1] [6] [12] [17] [25]. These approaches often improve forecasting accuracy or efficiency over plain GRU/LSTM baselines, especially in multiscale or irregular settings, but they do not eliminate long-horizon instability, error accumulation, or sensitivity to nonstationarity [18] [21] [70].

Relative to this literature, FAC-RNN cannot yet be directly positioned because the supplied evidence does not include a peer-reviewed paper explicitly describing FAC-RNN [31] [32] [33]. The most defensible conclusion is therefore comparative by omission: what appears missing or at least unverified for FAC-RNN is adaptive per-step timescale control, learned update-rate modulation, irregular-time state evolution, calibration analysis, and distribution-shift robustness testing [6] [12] [25] [36].

Overall confidence level: High - The field-level synthesis is strongly grounded across multiple themes, while the specific FAC-RNN comparison remains constrained by missing primary evidence rather than conflicting evidence [1] [12] [25] [31] [37].

References

1. Cui, J., Wang, X. (2026) AdaptiveFrequencyNet: a long-term time series forecasting approach based on learnable frequency decomposition and cross-scale attention fusion. *Scientific Reports*. <https://doi.org/10.1038/s41598-026-56619-x>
2. Yu, W., Zhou, J., Guo, W., (...), Luo, S. (2027) LMSNet: Variable-aware multi-scale framework for multivariate time series forecasting. *Neural Networks*. <https://doi.org/10.1016/j.neunet.2026.109333>
3. Wang, Q., Nicodemas, K.A. (2026) Hierarchical Attention Transformer for Multivariate Time Series Forecasting. *Computers, Materials and Continua*. <https://doi.org/10.32604/cmc.2026.074305>
4. Chen, D., Chen, L., Zhang, Y., (...), Yang, C. (2022) A Multiscale Interactive Recurrent Network for Time-Series Forecasting. *IEEE Transactions on Cybernetics*. <https://doi.org/10.1109/TCYB.2021.3055951>
5. Du, J., Jiang, S. (2026) HD-KAN: Hierarchical dual-dimension Kolmogorov-Arnold Networks with bidirectional time-frequency fusion for time series forecasting. *Information Sciences*. <https://doi.org/10.1016/j.ins.2026.123760>
6. Zhao, L., Li, F., Hu, G., (...), Xu, Y. (2026) MFformer: Time series forecasting via Fourier transform and meta-learning routing. *Computers and Electrical Engineering*. <https://doi.org/10.1016/j.compeleceng.2026.111142>
7. Xie, R., Zeng, A., Pan, D., Zhao, Y. (2026) CGFB-GNN: Explainable spatio-temporal GNN for cognition-guided functional brain network. *Biomedical Signal Processing and Control*. <https://doi.org/10.1016/j.bspc.2026.110842>
8. Messou, F.J.A., Chen, J., Zhang, S., (...), Yu, K. (2026) TAPformer : A transformer with adaptive attention and temporal projection encoding for long-term electricity load forecasting in IoT systems. *Internet of Things (The Netherlands)*. <https://doi.org/10.1016/j.iot.2026.101895>
9. Aslam, L., Zou, R., Huang, Y., (...), Yousaf, S. (2026) Physics-informed multi-gated convolutional recurrent network for extreme wind speed prediction. *Applied Energy*. <https://doi.org/10.1016/j.apenergy.2026.128094>
10. Yan, Fei, Jiang, Ying, Chen, Yuehao, Xu, Zhao-Dong (2025) A data-driven model for tunnel fire scenario deduction based on optimized informer. *Measurement Science and Technology*. <https://doi.org/10.1088/1361-6501/adcd8f>
11. Li, H., Wu, W., Chen, W., Zhang, M. (2025) DualNet: A dual-path network with adaptive compensation for multivariate time series forecasting. *Neurocomputing*. <https://doi.org/10.1016/j.neucom.2025.130720>
12. Tan, X., Wang, L., Wang, M., Zhang, Y. (2025) KOSLM: A Kalman-Optimal Hybrid State-Space Memory Network for Long-Term Time Series Forecasting. *Applied Sciences (Switzerland)*. <https://doi.org/10.3390/app152312684>
13. Li, Y. (2026) An Improved GRU Financial Time Series Prediction Model. *Fractal and Fractional*. <https://doi.org/10.3390/fractalfract10040227>

14. Ohno, K., Kumagai, A. (2021) Recurrent Neural Networks for Learning Long-term Temporal Dependencies with Reanalysis of Time Scale Representation. *Proceedings - 12th IEEE International Conference on Big Knowledge, ICBK 2021*. <https://doi.org/10.1109/ICKG52313.2021.00033>
15. Liu, J., Chen, S. (2019) Non-stationary Multivariate Time Series Prediction with Selective Recurrent Neural Networks. *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*. https://doi.org/10.1007/978-3-030-29894-4_51
16. Yadav, H., Thakkar, A. (2024) NOA-LSTM: An efficient LSTM cell architecture for time series forecasting. *Expert Systems with Applications*. <https://doi.org/10.1016/j.eswa.2023.122333>
17. Zhang, Q., Li, X., Wen, X., (...), Xu, Z. (2026) Servo motor forecasting in EV battery swapping systems via time-frequency and human-in-the-loop memory enhancement. *Journal of Manufacturing Systems*. <https://doi.org/10.1016/j.jmsy.2026.05.007>
18. Yu, B., Du, H., Zhang, D. (2026) Forecasting the realized volatility of oil futures: Machine learning with cross-market spillovers. *International Review of Economics and Finance*. <https://doi.org/10.1016/j.iref.2026.105176>
19. Li, S., Wu, Y., Fang, J., (...), Ling, J. (2026) Enhanced safety and emergency decision-making in NPPs: Multi-step parameter prediction for complex accident scenarios. *Progress in Nuclear Energy*. <https://doi.org/10.1016/j.pnucene.2025.106185>
20. Zhu, B., Connor, C. (2025) SugarNet: Personalized blood glucose forecast with a Fourier Kolmogorov–Arnold network. *Expert Systems with Applications*. <https://doi.org/10.1016/j.eswa.2025.127476>
21. Chen, X., Liu, T., Song, M., (...), Liu, Z. (2026) A Long-term forecasting model for sea surface temperature based on wavelet decomposition and adaptive fusion. *Journal of Sea Research*. <https://doi.org/10.1016/j.seares.2026.102696>
22. Qu, F., Wang, Z., Liu, Z., (...), Wang, L. (2026) A causality-inspired multi-dimensional feature fusion framework for multi-step engineering parameter prediction and reliability assessment. *Reliability Engineering and System Safety*. <https://doi.org/10.1016/j.ress.2026.112741>
23. Miao, H., Zhu, W., Dan, Y., Yu, N. (2024) Chaotic time series prediction based on multi-scale attention in a multi-agent environment. *Chaos, Solitons and Fractals*. <https://doi.org/10.1016/j.chaos.2024.114875>
24. Weerakody, P.B., Wong, K.W., Wang, G., Ela, W. (2021) A review of irregular time series data handling with gated recurrent neural networks. *Neurocomputing*. <https://doi.org/10.1016/j.neucom.2021.02.046>
25. Fitzgerald, O., Perez-Concha, O., Gallego-Luxan, B., (...), Jorm, L. (2023) Continuous time recurrent neural networks: Overview and benchmarking at forecasting blood glucose in the intensive care unit. *Journal of Biomedical Informatics*. <https://doi.org/10.1016/j.jbi.2023.104498>

26. Shi, Z., Wang, S., Yue, L., (...), Li, X. (2021) Deep dynamic imputation of clinical time series for mortality prediction. *Information Sciences*. <https://doi.org/10.1016/j.ins.2021.08.016>
27. Moualla, L., Rucci, A., Naletto, G., (...), Da Deppo, V. (2025) Hybrid GIS-Transformer Approach for Forecasting Sentinel-1 Displacement Time Series. *Remote Sensing*. <https://doi.org/10.3390/rs17142382>
28. Metzger, N., Türkoglu, M.O., D'Aronco, S., (...), Schindler, K. (2022) Crop Classification under Varying Cloud Cover with Neural Ordinary Differential Equations. *IEEE Transactions on Geoscience and Remote Sensing*. <https://doi.org/10.1109/TGRS.2021.3101965>
29. Fitzgerald, O., Perez-Concha, O., Gallego-Luxan, B., (...), Jorm, L. (2023) Continuous time recurrent neural networks: Overview and benchmarking at forecasting blood glucose in the intensive care unit. *Journal of Biomedical Informatics*. <https://doi.org/10.1016/j.jbi.2023.104498>
30. Chen, Yuming, Sanz-Alonso, Daniel, Willett, Rebecca (2023) Reduced-order autodifferentiable ensemble Kalman filters. *Inverse Problems*. <https://doi.org/10.1088/1361-6420/acff14>
31. Wang, T., Zhang, T., Zhang, K., (...), Lu, J. (2023) Context Attention Fusion Network for crowd counting. *Knowledge-Based Systems*. <https://doi.org/10.1016/j.knosys.2023.110541>
32. Zhu, Q., Jiang, X., Ye, R. (2021) Sentiment Analysis of Review Text Based on BiGRU-Attention and Hybrid CNN. *IEEE Access*. <https://doi.org/10.1109/ACCESS.2021.3118537>
33. Kamyab, M., Liu, G., Rasool, A., Adjeisah, M. (2022) ACR-SA: attention-based deep model through two-channel CNN and Bi-RNN for sentiment analysis. *PeerJ Computer Science*. <https://doi.org/10.7717/peerj-cs.877>
34. Zhang, W., Zhan, H., Sun, H., Yang, M. (2025) Probabilistic load forecasting for integrated energy systems based on quantile regression patch time series Transformer. *Energy Reports*. <https://doi.org/10.1016/j.egy.2024.11.057>
35. Li, Y., Wang, Z., Yuan, H., Bi, J. (2025) Hybrid water quality prediction based on attention combined with frequency enhancement and multi-seasonal decomposition. *Journal of Water Process Engineering*. <https://doi.org/10.1016/j.jwpe.2025.108747>
36. Ehmimed, Nadir, Chkouri, Mohamed Yassin, Touhafi, Abdellah (2025) Reliable and Adaptive Probabilistic Forecasting for Event-Driven Water-Quality Time Series Using a Gated Hybrid-Mixture Density Network. *Sensors (Basel, Switzerland)*. <https://doi.org/10.3390/s25247560>
37. Hewamalage, H., Bergmeir, C., Bandara, K. (2021) Recurrent Neural Networks for Time Series Forecasting: Current status and future directions. *International Journal of Forecasting*. <https://doi.org/10.1016/j.ijforecast.2020.06.008>
38. Lukoševičius, M., Jaeger, H. (2009) Reservoir computing approaches to recurrent neural network training. *Computer Science Review*. <https://doi.org/10.1016/j.cosrev.2009.03.005>
39. Huang, S.-C., Wu, C.-F., Chen, K.-C., (...), Chiou, C.-C. (2025) Intelligent dual-stage attention-based deep networks for energy market predictions. *Alexandria Engineering Journal*. <https://doi.org/10.1016/j.aej.2025.03.031>

40. Vlachas, P.R., Pathak, J., Hunt, B.R., (...), Koumoutsakos, P. (2020) Backpropagation algorithms and Reservoir Computing in Recurrent Neural Networks for the forecasting of complex spatiotemporal dynamics. *Neural Networks*. <https://doi.org/10.1016/j.neunet.2020.02.016>
41. Zhan, Y., Du, Z., Wu, B., (...), Gan, Y. (2026) Risk warning of transmission tower based on deformation monitoring system and enhanced RNN model. *Electric Power Systems Research*. <https://doi.org/10.1016/j.epsr.2026.112911>
42. Struye, J., Latré, S. (2020) Hierarchical temporal memory and recurrent neural networks for time series prediction: An empirical validation and reduction to multilayer perceptrons. *Neurocomputing*. <https://doi.org/10.1016/j.neucom.2018.09.098>
43. Lu, G., Ou, Y., Wang, Z., (...), Li, W. (2025) A Survey of Deep Learning for Time Series Forecasting: Theories, Datasets, and State-of-the-Art Techniques. *Computers, Materials and Continua*. <https://doi.org/10.32604/cmc.2025.068024>
44. Hao, X., Yin, H., Gao, J., Lan, H. (2025) Ultra short-term forecasting for the propulsion energy consumption of all-electric ships based on TCFFA-GRU-parallel network. *Results in Engineering*. <https://doi.org/10.1016/j.rineng.2024.103767>
45. Zhang, G.P. (2003) Time series forecasting using a hybrid ARIMA and neural network model. *Neurocomputing*. [https://doi.org/10.1016/S0925-2312\(01\)00702-0](https://doi.org/10.1016/S0925-2312(01)00702-0)
46. Lin, S., Lin, W., Wu, W., (...), Zhang, H. (2026) SegRNN: Segment Recurrent Neural Network for Long-Term Time-Series Forecasting. *IEEE Internet of Things Journal*. <https://doi.org/10.1109/JIOT.2025.3647705>
47. Fujita, T., Kawanishi, Y. (2027) Lightweight Gating Mechanism for RNNs with Sech-Based Vector Gates. *Lecture Notes in Computer Science*. https://doi.org/10.1007/978-3-032-31920-3_13
48. Faranda, D., Vrac, M., Yiou, P., (...), Gautard, V. (2021) Enhancing geophysical flow machine learning performance via scale separation. *Nonlinear Processes in Geophysics*. <https://doi.org/10.5194/npg-28-423-2021>
49. Sun, J., Sun, Z., Chen, Z., (...), Zhao, X. (2026) MSA-LR: Enhancing multi-scale temporal dynamics in multivariate time series forecasting with low-rank self-attention. *Neural Networks*. <https://doi.org/10.1016/j.neunet.2025.108150>
50. Shakhovska, N., Shymanskyi, V., Maherovskyi, A. (2026) A Methodology for Adaptive Design of Fractal Neural Architectures for Forecasting Self-Similar and Multifractal Time Series. *Applied Sciences (Switzerland)*. <https://doi.org/10.3390/app16157563>
51. Pourcel, G., Goldmann, M., Fischer, I., Soriano, M.C. (2024) Adaptive control of recurrent neural networks using conceptors. *Chaos*. <https://doi.org/10.1063/5.0211692>
52. Farooq, J., Abid Bazaz, M.A., Rafiq, D. (2023) Multiscale Autoencoder-RNN Architecture for Mitigating Error Accumulation in Long-Term Forecasting. *Proceedings of the 2023 International Conference on Emerging Techniques in Computational Intelligence, ICETCI 2023*. <https://doi.org/10.1109/ICETCI58599.2023.10331086>

53. Rodríguez-López, L., Usta, D.F.B., Bravo Alvarez, L.B., (...), Urrutia, R. (2026) Short-term forecasting of Chlorophyll-a dynamics in Lake Villarrica using foundation time-series models. *Ecological Informatics*. <https://doi.org/10.1016/j.ecoinf.2026.103938>
54. Jahani, A., Khalilpour, S.A., Seyedi, H., (...), Najafi, E. (2026) EP-RNN: Error-projected recurrent neural network for time series forecasting. *Results in Engineering*. <https://doi.org/10.1016/j.rineng.2026.110454>
55. Chen, S., Chen, K., Zhao, X., (...), Wang, W. (2026) ProtoPFL: Enhancing privacy and accuracy in federated time-series forecasting with heterogeneous data. *Computer Standards and Interfaces*. <https://doi.org/10.1016/j.csi.2026.104161>
56. Malin, M., Okkonen, J., Suutala, J. (2025) Snow water equivalent forecasting in sub-arctic and arctic regions: Efficient recurrent neural networks approach. *Environmental Modelling and Software*. <https://doi.org/10.1016/j.envsoft.2025.106695>
57. McHara, W., Raissi, M. (2026) An intelligent global irradiance forecasting framework based on a WOA-optimized FCM-CNN-WNN-Transformer hybrid network for solar electric vehicle energy management. *Measurement: Journal of the International Measurement Confederation*. <https://doi.org/10.1016/j.measurement.2026.121970>
58. Navidi, S., Bajarunas, K., Arias, M.A., Hu, C. (2025) Forecasting battery capacity for second-life applications using physics-informed recurrent neural networks. *eTransportation*. <https://doi.org/10.1016/j.etrans.2025.100432>
59. Teng, M., Meng, J., Pang, W., Wu, Y. (2027) Disentangled temporal modeling with sparse convolutions for time series forecasting. *Information Processing and Management*. <https://doi.org/10.1016/j.ipm.2026.104999>
60. Dong, C., Lei, L. (2026) A unified intelligent feature selection and adaptive learning framework for real-time cooling load forecasting and energy-efficient management in large commercial buildings. *Energy and Buildings*. <https://doi.org/10.1016/j.enbuild.2025.116804>
61. Joo, S., Lee, Y., Song, S.H., (...), Lee, J.I. (2025) Leveraging explainable AI for reliable prediction of nuclear power plant severe accident progression. *Reliability Engineering and System Safety*. <https://doi.org/10.1016/j.ress.2025.111307>
62. He, X., Shi, S., Geng, X., Xu, L. (2022) Information-aware attention dynamic synergetic network for multivariate time series long-term forecasting. *Neurocomputing*. <https://doi.org/10.1016/j.neucom.2022.04.124>
63. Yang, Yishuai (2025) TCN-QV: an attention-based deep learning method for long sequence time-series forecasting of gold prices. *PLoS ONE*. <https://doi.org/10.1371/journal.pone.0319776>
64. Orojo, O., Tepper, J., McGinnity, T.M., Mahmud, M. (2023) The Multi-Recurrent Neural Network for State-Of-The-Art Time-Series Processing. *Procedia Computer Science*. <https://doi.org/10.1016/j.procs.2023.08.187>

65. Yunita, A., Pratama, M.I., Almuzakki, M.Z., (...), Basori, A.H. (2025) Performance analysis of neural network architectures for time series forecasting: A comparative study of RNN, LSTM, GRU, and hybrid models. *MethodsX*. <https://doi.org/10.1016/j.mex.2025.103462>
66. Naghashi, V., Boukadoum, M., Diallo, A.B. (2025) Should We Reconsider RNNs for Time-Series Forecasting?. *AI (Switzerland)*. <https://doi.org/10.3390/ai6050090>
67. Khaldi, R., El Afia, A., Chiheb, R., Tabik, S. (2023) What is the best RNN-cell structure to forecast each time series behavior?. *Expert Systems with Applications*. <https://doi.org/10.1016/j.eswa.2022.119140>
68. Nawej, C.M., Owolawi, P.A., Walingo, T. (2025) Toward a Realistic Comparative Analysis of Recurrent Neural Network's Methods via Long-Term Memory Approaches. *Lecture Notes in Networks and Systems*. https://doi.org/10.1007/978-981-97-5035-1_19
69. Park, J., Park, C., Choi, J., Park, S. (2022) DeepGate: Global-local decomposition for multivariate time series modeling. *Information Sciences*. <https://doi.org/10.1016/j.ins.2022.01.015>
70. Wang, J., Ju, Y., Dong, P., Ju, T. (2025) Long-term time series forecasting by a frequency-domain enhanced temporal convolutional network with the stationary residual regularization. *Applied Soft Computing*. <https://doi.org/10.1016/j.asoc.2025.113779>